

1. An AI speaker is a typical use case for speech processing.

☒ True



☐ False

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing

2. Huawei's all-scenario AI solutions can be deployed in all kinds of environments, including public cloud, private clouds, various forms of edge computing, industrial IoT devices, and consumer devices.

☒ True



☐ False

Previous

Next



Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing

3. Typically, in a network that uses sigmoid, the gradient decreases to 0 within five layers, making the network difficult to train.

☒ True



☐ False

Previous

Next

✔ Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing



4. When you train a deep neural network with the softsign function, if the network is very deep, an increasing number of back-propagated gradients fall into the saturation area, making the gradient module smaller and finally close to 0, and the weight cannot be updated.

☒ True☐ False[Previous](#)[Next](#)

Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing



5. The tanh function can effectively solve the vanishing gradient problem.

☐ True

☒ False



Previous

Next

✓ Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing

6. L1 regularization is referred to as weight decay.

☐ True

☒ False



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing



7. If the hyperparameter search space is large, random search is better than grid search.

☒ True



☐ False

Previous

Next

✔ Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing



8. In a random search, each setting item is sampled from possible parameter values to find the optimal parameter subset.

☒ True



☐ False

Previous

Next

✔ Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing

9. Assuming a dataset contains the areas and prices of 21,613 housing units in a city, you can use a classification model to predict the prices of other housing units in the city.

☐ True

☒ False



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing

10. The emergence of foundation models represents a shift from weak AI to general AI.

☒ True



☐ False

Previous

Next

✖ Sorry, your answer was wrong. My Score: 0 Point(s) Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing



11. Batch normalization is a technique used in deep learning to optimize training. It reduces the dependence on initial weights and learning rates, resulting in improved model performance.

☒ True



☐ False

Previous

Next

✓ Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing

12. In regression tasks, the most commonly used loss function is cross entropy cost.

☐ True

☒ False



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing



13. In pipeline parallelism, the parameters of each layer of a model are split into multiple parts and distributed to multiple computing units. After the forward computation is complete, the gradients are combined on the master computing unit in the original splitting mode.

☐ True

☒ False



Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing



14. Model fine-tuning is to adjust the model structure and the number of trainable parameters based on the pre-trained model, and to use more diverse data for training.

☐ True

☒ False



Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: False

[Feedback on Answer](#)

Answer Explanations:

Nothing



15. The core idea of adapter tuning is to add adapter modules to specific layers of a pre-trained model. These modules usually consist of two small neural networks: one for dimensionality reduction and the other for augmentation.

☐ True

☒ False



Previous

Next

 Sorry, your answer was wrong. My Score: 0 Point(s) Answer: True

[Feedback on Answer](#)

Answer Explanations:

Nothing

16. Without considering any regularization terms, the support vectors of a support vector machine (SVM) are composed of ().

- ☐ A. Points on the separating hyperplane
- ☐ B. The points farthest from the separating hyperplane
- ☒ C. The points closest to the separating hyperplane
- ☐ D. Points of a certain type

[Previous](#)[Next](#)

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: C

[Feedback on Answer](#)

Answer Explanations:

Nothing



17. During neural network training, which of the following values is continuously updated using the gradient descent method to minimize the loss function?

- ☐ A. Hyperparameters
- ☐ B. Feature value
- ☐ C. Number of samples
- ☒ D. Parameters



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: D

[Feedback on Answer](#)

Answer Explanations:

Nothing



18. The ReLU function is commonly used in deep learning neural networks. Which of the following is the value range of this function?

☒ A. $[0, +\infty)$



☐ B. $[0, 1]$

☐ C. $[-1, 1]$

☐ D. $[-1, 0]$

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: A

[Feedback on Answer](#)

Answer Explanations:

Nothing

19. The sigmoid activation function is monotonic and continuous, has bounded outputs, and makes the network easy to converge. It was popular for a while, but when the network is deep, which one of these problems is sigmoid associated with?

☐ A. Gradient reduction

☒ B. Vanishing gradient



☐ C. XOR Problem

☐ D. Overfitting

Previous

Next



Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: B

[Feedback on Answer](#)

Answer Explanations:

Nothing

20. Which of the following functions can be used to alleviate the vanishing gradient problem?

- ☐ A. Sigmoid
- ☐ B. Tanh
- ☐ C. Softsign
- ☒ D. ReLU



Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: D

[Feedback on Answer](#)

Answer Explanations:

Nothing



21. If a 32 x 32 image is input and a 5 x 5 kernel is used for convolution with the stride of 1, the size of the output image is ().

☐ A. 28 x 28

☒ B. 29 x 29



☐ C. 28 x 23

☐ D. 23 x 23

Previous

Next

Sorry, your answer was wrong. My Score: 0 Point(s) Answer: A

[Feedback on Answer](#)

Answer Explanations:

Nothing



22. A manufacturer wants to produce virtual assistants for hospital use. For virtual assistants, which of the following technologies is used for voiceprint recognition?

- ☒ A. Speech recognition and processing
- ☐ B. Image recognition and processing
- ☐ C. Expert system and knowledge graph
- ☐ D. Image generation and enhancement



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: A

[Feedback on Answer](#)

Answer Explanations:

Nothing



23. Which of the following is NOT a use case for natural language processing?

- ☐ A. Public opinion analysis
- ☐ B. Machine translation
- ☐ C. Text classification
- ☒ D. Image recognition



Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: D

[Feedback on Answer](#)

Answer Explanations:

Nothing

24. Which of the following is NOT a recurrent neural network?

☐ A. RNN

☐ B. LSTM

☒ C. CNN



☐ D. RNN

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: C

[Feedback on Answer](#)

Answer Explanations:

Nothing



25. Within the broad landscape of large language models built on the Transformer architecture, which roadmap does GPT belong to?

- ☐ A. Encoder-Only
- ☒ B. Decoder-Only
- ☐ C. Encoder-Decoder
- ☐ D. Decoder-Encoder



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: B

[Feedback on Answer](#)

Answer Explanations:

Nothing



26. A computer uses labeled images to learn and distinguish between images of apples and those of pears. What is the name of this learning method?

- ☒ A. Supervised learning
- ☐ B. Unsupervised learning
- ☐ C. Semi-supervised learning
- ☐ D. Reinforcement learning



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: A

[Feedback on Answer](#)

Answer Explanations:

Nothing



27. Which of the following statements is true about classification models and regression models in machine learning?

- ☐ A. For regression problems, the output variables are discrete values. For classification problems, the output variables are continuous.
- ☐ B. The most commonly used indicators for evaluating regression and classification problems are the accuracy and the recall rate.
- ☒ C. Overfitting may occur in both regression and classification problems. 
- ☐ D. Logistic regression is a typical regression model.

Previous

Next

 Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: C

[Feedback on Answer](#)

Answer Explanations:

Nothing



28. Which of the following statements about support vector machines (SVMs) is false?

- ☐ A. SVMs are classification models. Their basic model is the linear classifier that maximizes the width of the gap between the two categories in the feature space.
- ☐ B. SVMs also have a kernel trick, which makes them non-linear classifiers.
- ☐ C. In the case of linear inseparability, non-linear mapping algorithms are used to convert the linearly inseparable samples of low-dimensional input space into samples of high-dimensional feature space. In this way, samples become linearly separable.
- ☒ D. SVMs only apply to linear classification. 

Previous

Next



Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: D

[Feedback on Answer](#)

Answer Explanations:

Nothing

29. Which of the following is NOT an AI deep learning framework?

☐ A. TensorFlow

☒ B. Scikit-learn



☐ C. Pytorch

☐ D. Caffe

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: B

[Feedback on Answer](#)

Answer Explanations:

Nothing

30. Which of the following are used in TensorFlow to describe the computation process?

☐ A. Parameter

☐ B. Session

☒ C. Data flow



☐ D. Tensor

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: C

[Feedback on Answer](#)

Answer Explanations:

Nothing

31. If a 32×32 feature map is processed through a pooling layer with a stride of 2, the size of the output map is ().

☐ A. 30×30

☒ B. 16×16



☐ C. 29×29

☐ D. 23×23

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: B

[Feedback on Answer](#)

Answer Explanations:

Nothing



32. Which of the following is NOT a characteristic of PyTorch?

- ☐ A. Python first
- ☐ B. Easy to debug
- ☐ C. Extensive ecosystem
- ☒ D. Default use of static maps



Previous

Next

✓ Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: D

[Feedback on Answer](#)

Answer Explanations:

Nothing



33. Which of the following is not an advanced AI application?

- ☐ A. Knowledge graph
- ☒ B. rget
Quantum computing
- ☐ C. SDN
- ☐ D. Intelligent driving



Previous

Next

 Sorry, your answer was wrong. My Score: 0 Point(s) Answer: C

[Feedback on Answer](#)

Answer Explanations:

Nothing



34. Which of the following statements is true about the training of large models?

- ☐ A. Foundation models cannot be directly deployed.
- ☒ B. The pre-training of large models requires a large amount of high-quality data. ✓
- ☐ C. Unlabeled data used in the pre-training phase of large models must be used in the fine-tuning phase.
- ☐ D. Large models have emergent abilities, and therefore they do not need to be aligned with human values.

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s) Answer: B

[Feedback on Answer](#)

Answer Explanations:

Nothing

35. Tensor parallelism is a common method for parallel training of large models. Which of the following types of model data is split to different devices?

- ☐ A. Gradient
- ☐ B. Optimizer
- ☐ C. Input data
- ☒ D. Model parameters



Previous

Next



Congratulations! Your answer was correct.

My Score: 16 Point(s)

Answer: D

[Feedback on Answer](#)

Answer Explanations:

Nothing



36. What are the layers used in Transformer to capture time dependencies within a sequence?

☐ A. Residual layer

☒ B. Self-attention layer



☒ C. Location code



☐ D. Feed-forward network layer

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: BC

[Feedback on Answer](#)

Answer Explanations:

Nothing

37. During the calculation of the attention mechanism, which of the following are required for similarity calculation to obtain the attention score?

- ☒ A. Query ✓
- ☐ B. Weight
- ☐ C. Value
- ☒ D. Key ✓

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: AD

[Feedback on Answer](#)

Answer Explanations:

Nothing

38. Which of the following are common evaluation metrics for classification tasks?

- ☒ A. Accuracy ✓
- ☒ B. Precision ✓
- ☒ C. Recall rate ✓
- ☐ D. Mean squared error

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABC

[Feedback on Answer](#)

Answer Explanations:

Nothing

39. Which of the following hyperparameters may be defined during model training?

- ☒ A. Batch size ✓
- ☒ B. Learning rate ✓
- ☒ C. Optimizer type ✓
- ☐ D. Loss value

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABC

[Feedback on Answer](#)

Answer Explanations:

Nothing



40. Optimizers are often used to accelerate training of deep learning neural networks. Which of the following are deep learning optimizers?

☒ A. Momentum



☒ B. Adagrad



☒ C. RMSprop



☒ D. Adam



Previous

Next

 Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABCD

[Feedback on Answer](#)

Answer Explanations:

Nothing



41. Which of the following statements about some of the subfields of AI are true?

- ☒ A. Computer vision is a science that studies how to make computers "see" things. ✓
- ☒ B. Speech processing is a general term for different speech processing methods and techniques, including the study of vocalization, the collection of statistics related to speech signals, speech recognition, speech synthesis, and speech perception. ✓
- ☒ C. Natural language processing studies how to use computer technology to understand and use human language. ✓
- ☐ D. Autonomous driving does not require speech processing or computer vision.

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABC

[Feedback on Answer](#)

Answer Explanations:

Nothing

42. Which of the following are common activation functions of neural networks?

☐ A. Dropout

☒ B. Sigmoid



☒ C. tanh



☒ D. Leaky ReLU



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: BCD

[Feedback on Answer](#)

Answer Explanations:

Nothing

43. Which of the following comprise the Adam optimizer?

☒ A. Momentum



☐ B. Adagrad

☒ C. RMSprop



☐ D. Nesterov

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: AC

[Feedback on Answer](#)

Answer Explanations:

Nothing

44. Which of the following statements about datasets are true?

- ☒ A. A dataset is a collection of data used in machine learning tasks. Each piece of data is called a sample. ✓
- ☒ B. Events or attributes that reflect the particular performance or nature of a sample are called features. ✓
- ☒ C. A training set is a dataset used in the training process, where each sample is referred to as a training sample. ✓
- ☒ D. Learning (or training) is the process of building a model from data. ✓

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABCD

[Feedback on Answer](#)

Answer Explanations:

Nothing



45. Which of the following statements about data preprocessing are true?

- ☒ A. Data cleansing is a process of filling in missing values, as well as detecting and eliminating noise and exceptions. ✓
- ☒ B. Data dimension reduction aims to simplify data attributes and avoid the curse of dimensionality. ✓
- ☒ C. Data standardization aims to reduce noise data and improve model accuracy by standardizing data. ✓
- ☐ D. Machine learning tasks generate results in the form of model outputs. Therefore, model training is more important than data preprocessing.

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABC

[Feedback on Answer](#)

Answer Explanations:

Nothing



46. Which of the following statements about model parameters and hyperparameters are true?

☒ A. Models contain both parameters and hyperparameters.



☐ B. Hyperparameters are automatically learned by models.

☒ C. Hyperparameters are manually set.



☒ D. Hyperparameters can be used to control training.



Previous

Next

 Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ACD

[Feedback on Answer](#)

Answer Explanations:

Nothing



47. Data augmentation can improve model robustness and avoid overfitting. Which of the following are data augmentation methods?

☐ A. Batch normalization

☒ B. Image contrast change



☒ C. Image noise adding



☒ D. Random crop



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: BCD

[Feedback on Answer](#)

Answer Explanations:

Nothing



48. The softmax function is frequently used in classification tasks. Which of the following statements about it are true?

- ☒ A. The softmax function is also called the normalized exponential function. ✓
- ☐ B. The softmax regression model is an algorithm for solving binary classification regression problems.
- ☒ C. The softmax function is the generalization of the binary classification function sigmoid. ✓
- ☒ D. The softmax function is usually used together with the cross-entropy loss function. ✓

[Previous](#)[Next](#)

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ACD

[Feedback on Answer](#)

Answer Explanations:

Nothing

49. Which of the following are characteristics of eager execution in TensorFlow 2.x?

☐ A. High performance

☐ B. Deployability

☒ C. Intuitiveness



☒ D. Flexibility



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: CD

[Feedback on Answer](#)

Answer Explanations:

Nothing

50. The process of developing AI applications includes ().

- ☒ A. Data preprocessing ✓
- ☒ B. Modeling ✓
- ☒ C. Model training ✓
- ☒ D. Model deployment ✓

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABCD

[Feedback on Answer](#)

Answer Explanations:

Nothing



51. Which of the following statements about large models are true?

- ☒ A. A large amount of data is required for training. ✓
- ☒ B. The model parameter scale is large. ✓
- ☐ C. Training data does not need to be cleansed because this helps models learn how to process error information.
- ☒ D. The generalization capability of large models is stronger than that of small models. ✓

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABD

[Feedback on Answer](#)

Answer Explanations:

Nothing

52. Which of the following fields can large model technologies be applied to?

- ☒ A. Autonomous driving ✓
- ☒ B. Robots ✓
- ☒ C. Translation ✓
- ☒ D. Finance ✓

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: ABCD

[Feedback on Answer](#)

Answer Explanations:

Nothing

53. Which of the following may be included in the training dataset of GPT-3?

☐ A. COCO2017

☒ B. Common Crawl



☒ C. Wikipedia



☒ D. Books1



Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: BCD

[Feedback on Answer](#)

Answer Explanations:

Nothing

54. Which of the following fine-tuning methods do not change the Transformer block structure of a large model?

☐ A. Adapter tuning

☒ B. Prefix tuning



☒ C. Prompt tuning



☐ D. LoRA

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 18 Point(s) Answer: BC

[Feedback on Answer](#)

Answer Explanations:

Nothing



55. The development of large models requires not only training and fine-tuning, but also evaluation. Which of the following are reasons for evaluating large models?

- ☒ A. Improving the accuracy of model outputs ✗
- ☒ B. Evaluating the capabilities of large models ✓
- ☒ C. Providing reference for enterprises to select models ✓
- ☒ D. Aligning large models with human values ✗

[Previous](#)[Next](#)

✗ Sorry, your answer was wrong. My Score: 0 Point(s) Answer: BC

[Feedback on Answer](#)

Answer Explanations:

Nothing

56. The three essential elements for AI development are data, compute, and ____.

1 Algorithm



Answer: Algorithm

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s)

[Feedback on Answer](#)

Answer Explanations:

Nothing

57. CNN stands for ____.

1 Convolutional Neural Network



Answer: Convolutional neural network

Previous

Next

✔ Congratulations! Your answer was correct. My Score: 16 Point(s)

[Feedback on Answer](#)

Answer Explanations:

Nothing



58. A/An ____ is a binary classification model that is based on a linear classifier defined in the feature space. (Enter an abbreviation.)

1 SVM



Answer: SVM

Previous

Next

✓ Congratulations! Your answer was correct. My Score: 16 Point(s)

[Feedback on Answer](#)

Answer Explanations:

Nothing

59. The ____ algorithm is based on the idea that if most of a sample's k nearest neighbors in the feature space belong to the same class, the sample also belongs to that class. (Enter an abbreviation.)

1 KNN



Answer: KNN

[Previous](#)[Next](#)

✔ Congratulations! Your answer was correct. My Score: 16 Point(s)

[Feedback on Answer](#)**Answer Explanations:**

Nothing

60. In k-fold cross-validation, the existing data is divided into ____ groups.

1 K ✓

Answer: k

Previous

✓ Congratulations! Your answer was correct. My Score: 16 Point(s)

[Feedback on Answer](#)

Answer Explanations:

Nothing